

Project 3 Showcase

We'll be running a Project 3 Showcase day Sunday, May 11, from 12:00-15:00 in the Woz.

Show off your project 3 to other students! Or come by and see what creative ideas were developed!

Registration link: <https://forms.gle/UmA8YkNBqm8YBa8A>



Lecture 37 (Sorting 6)

Radix vs. Comparison Sorting

CS61B, Spring 2025 @ UC Berkeley

Slides credit: Josh Hug

Radix Sorting Strings

Lecture 37, CS61B, Spring 2025

Sorting Conclusion

- **Radix Sorting Strings**
- Radix Sorting Integers
- Sound of Sorting

Algorithm Design Practice

- Abstracting Data Structures

Practice Problems

Merge Sort requires $\Theta(N \log N)$ compares.

What is Merge Sort's runtime on strings of length W ? (Are comparisons necessarily constant time anymore?)

Merge Sort requires $\Theta(N \log N)$ compares.

What is Merge Sort's runtime on strings of length W ?

- It depends!
 - $\Theta(N \log N)$ if each comparison takes constant time.
 - Example: Strings are all different in top character.
 - $\Theta(WN \log N)$ if each comparison takes $\Theta(W)$ time.
 - Example: Strings are all equal.

The facts.

- Treating alphabet size as constant, LSD Sort has runtime $\Theta(WN)$.
- Merge Sort has runtime between $\Theta(N \log N)$ and $\Theta(WN \log N)$.

Which is better? It depends.

- When might LSD sort be faster?
- When might Merge Sort be faster?

The facts:

- Treating alphabet size as constant, LSD Sort has runtime $\Theta(WN)$.
- Merge Sort is between $\Theta(N \log N)$ and $\Theta(WN \log N)$.

Which is better? It depends.

- When might LSD sort be faster?
 - Sufficiently large N .
 - If strings are very similar to each other.
 - Each Merge Sort comparison costs $\Theta(W)$ time.
- When might Merge Sort be faster?
 - If strings are highly dissimilar from each other.
 - Each Merge Sort comparison is very fast.

AAAAAAAAAAAAA.....AB
AAAAAAAAAAAAA.....AA
AAAAAAAAAAAAA.....AQ
...
IUYQWLKJASHLEIUHAD...
LIUHLIUHRGLIUHWEF...
OZIUHIOHLHLZIEIUHF...
...

Radix Sorting Integers

Lecture 37, CS61B, Spring 2025

Sorting Conclusion

- Radix Sorting Strings
- **Radix Sorting Integers**
- Sound of Sorting

Algorithm Design Practice

- Abstracting Data Structures

Practice Problems

Issue: We don't have a `charAt` method for integers.

- How would you LSD radix sort an array of integers?

Issue: We don't have a `charAt` method for integers.

- How would you LSD radix sort an array of integers?
 - Could convert into a `String` and treat as a base 10 number. Since maximum Java int is 2,000,000,000, W is also 10.
 - Could modify LSD radix sort to work natively on integers.
 - Instead of using `charAt`, maybe write a helper method like `getDthDigit(int N, int d)`. Example: `getDthDigit(15009, 2) = 5`.

LSD Radix Sort on Integers

Note: There's no reason to stick with base 10!

- Could instead treat as a base 16, base 256, base 65536 number.

Example: 512,312 in base 16 is a 5 digit number:

- $512312_{10} = (7 \times 16^4) + (13 \times 16^3) + (1 \times 16^2) + (3 \times 16^1) + (8 \times 16^0)$

Note this digit is greater than 9! That's OK, because we're in base 16.

Example: 512,312 in base 256 is a 3 digit number:

- $512312_{10} = (7 \times 256^2) + (209 \times 256^1) + (56 \times 256^0)$

Note these digit are greater than 9! That's OK, because we're in base 256.

For Java integers:

- $R=10$, treat as a base 10 number. Up to 10 digits.
- $R=16$, treat as a base 16 number. Up to 8 digits.
- $R=256$, treat as a base 256 number. Up to 4 digits.
- $R=65536$, treat as a base 65536 number. Up to 2 digits.
- $R=2147483647$, treat as a base 2147483647 number (this is equivalent to counting sort). Has exactly 1 digit.

Interesting fact: Runtime depends on the alphabet size.

- As we saw with city sorting last time, $R = 2147483647$ will result in a very slow radix sort (since it's just counting sort).

Results of a computational experiment:

- Treating as a base 256 number (4 digits), LSD radix sorting integers easily defeats Quicksort.

Sort	Base	# of Digits	Runtime
Java QuickSort	N/A	N/A	10.9 seconds
LSD Radix Sort	$2^4 = 16$	8	3.6 seconds
LSD Radix Sort	$2^8 = 256$	4	2.28 seconds
LSD Radix Sort	$2^{16} = 65536$	2	3.66 seconds
LSD Radix Sort	$2^{30} = 1073741824$	2	20 seconds

Sorting 100,000,000 integers

Sorting Summary

Lecture 37, CS61B, Spring 2025

Sorting Conclusion

- Radix Sorting Strings
- Radix Sorting Integers
- **Sorting Summary**

Algorithm Design Practice

- Abstracting Data Structures

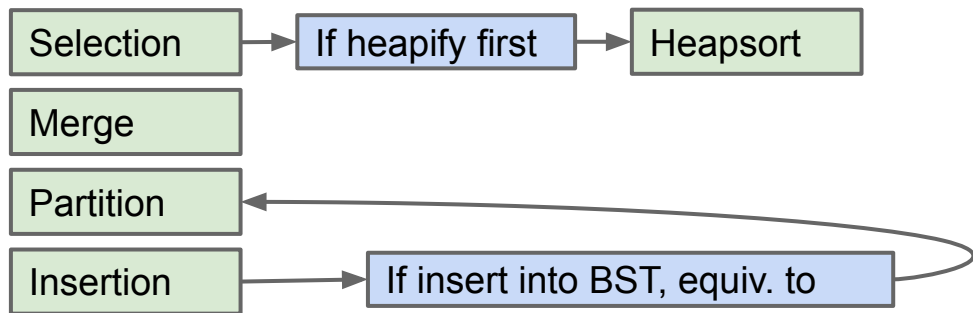
Practice Problems

Sorting Landscape


Below, we see the landscape of the sorting algorithms we've studied.

- Three basic flavors: Comparison, Alphabet, and Radix based.
- Each can be useful in different circumstances, but the important part was the analysis and the deep thought!
 - Hoping to teach you how to approach problems in general.

Comparison Based Sorting Algorithms:




Small-Alphabet (e.g. Integer) Sorting Algorithms:



Radix Sorting Algorithms:

(require a sorting subroutine)



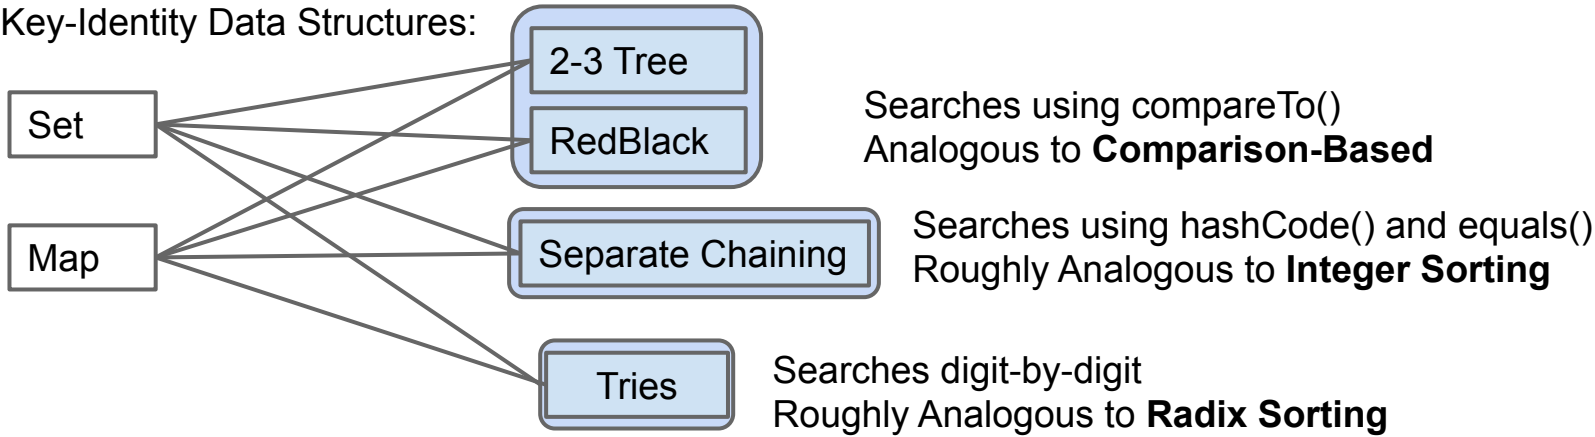
We've now concluded our study of the “sort problem.”

- During the data structures part of the class, we studied what we called the “search problem”: Retrieve data of interest.
- There are some interesting connections between the two.


Name	Storage Operation(s)	Primary Retrieval Operation	Retrieve By:
List	<code>add(key)</code> <code>insert(key, index)</code>	<code>get(index)</code>	index
Map	<code>put(key, value)</code>	<code>get(key)</code>	key identity
Set	<code>add(key)</code>	<code>containsKey(key)</code>	key identity
PQ	<code>add(key)</code>	<code>getSmallest()</code>	key order (a.k.a. key size)
Disjoint Sets	<code>connect(int1, int2)</code>	<code>isConnected(int1, int2)</code>	two int values

Partial list of search problem data structures.


Search-By-Key-Identity Data Structures:




Comparison Based Sorting Algorithms:



Radix Sorting Algorithms:
(require a sorting subroutine)



Small-Alphabet (e.g. Integer) Sorting Algorithms:



There's plenty more to explore!

Many of these ideas can be mixed and matched with others. Examples:

- What if we use quicksort as a subroutine for MSD radix sort instead of counting sort?
- Implementing the `comparable` interface means an object can be stored in our `compareTo`-based data structures (e.g. `TreeSet`), or sorted with our comparison based sorts. Is there a single equivalent interface that would allow storage in a trie AND radix sorting? What would that interface look like?
- If an object has both digits AND is comparable, could we somehow use an LLRB to improve radix sort in some way?

Sounds of Sorting Algorithms

Starts with selection sort: <https://www.youtube.com/watch?v=kPRA0W1kECg>

Insertion sort: <https://www.youtube.com/watch?v=kPRA0W1kECg&t=0m9s>

Quicksort: <https://www.youtube.com/watch?v=kPRA0W1kECg&t=0m38s>

Mergesort: <https://www.youtube.com/watch?v=kPRA0W1kECg&t=1m05s>

Heapsort: <https://www.youtube.com/watch?v=kPRA0W1kECg&t=1m28s>

LSD sort: <https://www.youtube.com/watch?v=kPRA0W1kECg&t=1m54s>

MSD sort: <https://www.youtube.com/watch?v=kPRA0W1kECg&t=2m10s>

Shell's sort: <https://www.youtube.com/watch?v=kPRA0W1kECg&t=3m37s>

More sorts: https://www.youtube.com/watch?v=8MsTNqK3o_w

Questions to ponder (later... after class):

- How many items are sorted in the video for selection sort?
- Why does insertion sort take longer / more compares than selection sort?
- At what time stamp does the first partition complete for Quicksort?
- Could the size of the input used by mergesort in the video be a power of 2?
- What do the colors mean for heapsort?
- How many characters are in the alphabet used for the LSD sort problem?
- How many digits are in the keys used for the LSD sort problem?

Abstracting Data Structures

Lecture 37, CS61B, Spring 2025

Sorting Conclusion

- Radix Sorting Strings
- Radix Sorting Integers
- Sorting Summary

Algorithm Design Practice

- **Abstracting Data Structures**

Practice Problems

- Lists/Deque
 - Store a collection of items in order.
 - Adding/removing from end in constant (amortized) time
 - Getting/Setting in constant time
 - Sorting in $N \log N$ time (Quicksort, Mergesort, etc.)
- Disjoint Set
 - Store a graph
 - Can determine if two nodes are connected in $\alpha(N)$ time
- Set/Map
 - Store a collection of items with no order (or key-value pairs)
 - Adding/removing any item in constant (amortized) time
 - Getting/setting in constant (amortized) time
- Heap
 - Store a collection of items
 - Adding/removing the smallest in $\log(N)$ time
- Graph
 - Store a graph
 - Dijkstra's, A*, Prim's, Kruskal's, etc.

Problems we've solved

- Given a collection of N items with an order, we can
 - Sort them in $N \log N$ time
 - Add to the collection in constant time
 - Iterate over the items with constant overhead (constant time per item, N time total)
 - Repeatedly find the smallest in $\log N$ time (if we make our adds a bit slower)
- Given a collection of N distinct items with no order, we can
 - Add to/remove from the collection in constant time
 - Iterate over the items with constant overhead
 - Associate a value to each item (with no additional overhead)
 - Check if a particular item is in there in constant time
- Given a graph, we can
 - Find the shortest path from between two vertices in $E \log V$ time
 - Find the minimum spanning tree in $E \log V$ time
 - Find whether two nodes are connected in $\alpha(N)$ time

Practice Problems

Lecture 37, CS61B, Spring 2025

Sorting Conclusion

- Radix Sorting Strings
- Radix Sorting Integers
- Sorting Summary

Algorithm Design Practice

- Abstracting Data Structures

Practice Problems

How to solve Algorithm Design Problems

Algorithm Design and Asymptotic Analysis are the two most "creative" parts of 61B

- Large portion of CS theory
- First step when making a new algorithm (don't start coding until you have a plan!)
- Often tested in interviews

There's no easy "trick" to solving these; the easiest way to get better is through practice.

Fortunately, there are many programming challenge websites which you can use to practice.

- Leetcode, Codeforces: Similar to interview questions, harder ones veer into competitive programming territory
- Advent of Code: Yearly code challenge that releases one problem per day in December
- Project Euler: More mathy than the others, less focused on runtime bounds

The rest of today will be algorithm design practice (goal is to describe an algorithm within a given runtime bound)

For each problem, I'll include the approximate difficulty of coming up with the algorithm from scratch (3/5 is around the difficulty of the hardest problems we'll ask on final exams). No specific information/algorithm in this section is considered in scope.

Duplicate

You are given a list of N integers (unsorted). Determine if the list contains at least one pair of duplicates

Example:

[1,2,3,10,5,8,10] -> True

[1,4,0,-5] -> False

Runtime Requirements:

$O(N \log N)$: ★☆☆☆☆

$O(N)$: ★★☆☆☆

Duplicate

You are given a list of N integers (unsorted). Determine if the list contains at least one pair of duplicates

Example:

[1,2,3,10,5,8,10] -> True

[1,4,0,-5] -> False

Runtime Requirements:

$O(N \log N)$: ★☆☆☆☆

Sort the list, then run dup2 from Lecture 15


$O(N)$: ★★☆☆☆

Insert all integers into a set, then check if that set contains exactly N items

You are given a graph whose edges are values between 0 and 1 (not inclusive), a start vertex, and an end vertex. Find the path whose edges **multiply to the largest value**.

Example: $0.6 * 0.7 * 0.8 = 0.336$, which is the largest possible value attainable in the below graph

Difficulty: ★★☆☆



You are given a graph whose edges are values between 0 and 1 (not inclusive), a start vertex, and an end vertex. Find the path whose edges **multiply to the largest value**.

Example: $0.6 * 0.7 * 0.8 = 0.336$, which is the largest possible value attainable in the below graph

Difficulty: ★★☆☆☆

Solution 1: Run Dijkstra's algorithm, but multiply instead of adding the weights, and use a max heap instead of a min heap. This works because the product of weights always decreases (you can't take an edge and increase the product).

You are given a graph whose edges are values between 0 and 1 (not inclusive), a start vertex, and an end vertex. Find the path whose edges **multiply to the largest value**.

Example: $0.6 * 0.7 * 0.8 = 0.336$, which is the largest possible value attainable in the below graph

Difficulty: ★★☆☆☆

Solution 1: Run Dijkstra's algorithm, but multiply instead of adding the weights, and use a max heap instead of a min heap. This works because the product of weights always decreases (you can't take an edge and increase the product).

Solution 2: (★★★★☆) Create a new graph, but replace each edge of weight w with an edge of weight $-\log_2(w)$ (this is guaranteed to be positive, since the edge weights are between 0 and 1).

Minimizing the sum of $-\log(x) - \log(y) - \log(z) - \dots$

is the same as maximizing the sum of $\log(x) + \log(y) + \log(z) + \dots$,

which is the same as maximizing $\log(xyz\dots)$,

which is the same as maximizing $xyz\dots$


So if we run Dijkstra's on the new graph, we'll get the correct path.

A Slight Modification

You are given a graph whose edges are values between 0 and 1 (not inclusive), a start vertex, and an end vertex. Find the path whose edges **add to the largest value without visiting any node more than once**.

Example: $0.6 + 0.7 + 0.8 = 0.21$, which is the largest possible value attainable in the below graph

Difficulty: ★★★★★



A Slight Modification

You are given a graph whose edges are values between 0 and 1 (not inclusive), a start vertex, and an end vertex. Find the path whose edges **add to the largest value without visiting any node more than once**.

Example: $0.6 + 0.7 + 0.8 = 0.21$, which is the largest possible value attainable in the below graph

Difficulty: ★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★★

Currently no known solution in polynomial time (and if you do solve this in polynomial time, you earn \$1000000)

Just like in asymptotic analysis, it's often hard to determine the difference between an easy and a hard problem.

Why do you earn \$1,000,000 for solving this problem?

Reducing Hamilton Path

The **Hamilton Path Problem** is the problem of, given a graph, determining if there is a path that goes through all vertices exactly once.

Problem X is the problem of, given a graph whose edges are values between 0 and 1 (not inclusive), a start vertex, and an end vertex, finding the path whose edges add to the largest value without visiting any node more than once.

Show that the Hamilton Path Problem **reduces** to Problem X; that is, if you find an efficient solution to Problem X, then you have an efficient solution to the Hamilton Path Problem (★★★☆☆)

Reducing Hamilton Path

The **Hamilton Path Problem** is the problem of, given a graph, determining if there is a path that goes through all vertices exactly once.

Problem X is the problem of, given a graph whose edges are values between 0 and 1 (not inclusive), a start vertex, and an end vertex, finding the path whose edges add to the largest value without visiting any node more than once.

Show that the Hamilton Path Problem **reduces** to Problem X; that is, if you find an efficient solution to Problem X, then you have an efficient solution to the Hamilton Path Problem (★★★☆☆)

Create a new graph with the same topology as the original graph, but with edge weights 0.5. Add two new "dummy" nodes S and T that are connected to all other nodes with weight 0.1. Run Problem X on this new graph from S to T, and return True if the path returned has length $0.2 + 0.5 \times (|V| - 1)$

Beans and Plates

A line of W plates are placed in a row, and numbered $1, 2, \dots, W$.

You're given N pairs of numbers (a,b) , $1 \leq a \leq b \leq W$.

For each pair of numbers, you place one bean on each plate from plate a to plate b (inclusive).

Example:

If $W = 5$, and you're given the pairs $(2,5)$, $(1,4)$, $(3,3)$, the plates look like:

$[0,0,0,0,0] \rightarrow [0,1,1,1,1] \rightarrow [1,2,2,2,1] \rightarrow [1,2,3,2,1]$

Find the total number of beans on all plates in $O(NW)$ time (☆☆☆☆☆)

Find the total number of beans on all plates in $O(N)$ time (★★★★☆)

Find the number of plates with an odd number of beans on them in $O(N \log N)$ time (★★★★★)

Beans and Plates

A line of W plates are placed in a row, and numbered $1, 2, \dots, W$.

You're given N pairs of numbers (a,b) , $1 \leq a \leq b \leq W$.

For each pair of numbers, you place one bean on each plate from plate a to plate b (inclusive).

Example:

If $W = 5$, and you're given the pairs $(2,5)$, $(1,4)$, $(3,3)$, the plates look like:

$[0,0,0,0,0] \rightarrow [0,1,1,1,1] \rightarrow [1,2,2,2,1] \rightarrow [1,2,3,2,1]$

Find the total number of beans on all plates in $O(NW)$ time (☆☆☆☆☆)

Simulate the procedure on an array of length W

Find the total number of beans on all plates in $O(N)$ time (★★★★☆)

Each pair adds $b-a+1$ beans in total. Sum this over all pairs

Beans and Plates

A line of W plates are placed in a row, and numbered $1, 2, \dots, W$.

You're given N pairs of numbers (a,b) , $1 \leq a \leq b \leq W$.

For each pair of numbers, you place one bean on each plate from plate a to plate b (inclusive).

Example:

If $W = 5$, and you're given the pairs $(2,5)$, $(1,4)$, $(3,3)$, the plates look like:

$[0,0,0,0,0] \rightarrow [0,1,1,1,1] \rightarrow [1,2,2,2,1] \rightarrow [1,2,3,2,1]$

Find the number of plates with an odd number of beans on them in $O(N \log N)$ time (★★★★★)

Create a list `lst`, and insert into the list a and $b+1$ for each pair (In the example, we insert $2,6,1,5,3,4$).

Sort this list. This creates a list of every time the number of beans in the list changes from odd to even or vice versa. So the odd-bean plates are the ones between `lst[0]` and `lst[1]-1`, `lst[2]` and `lst[3]-1`, etc. Sum these values together.